

Local Search and Optimization Problems

Dr. Sandareka Wickramanayake

Outline

- Hill Climbing
- Simulated Annealing
- Local Beam Search
- Genetic Algorithms

Reference

- Artificial Intelligence - A Modern Approach – Chapter 4 – Section 1

Search Types

- Backtracking state-space search
- **Local Search and Optimization**
- Constraint satisfaction search
- Adversarial search

Local Search and Optimization

- Previous searches: **keep paths in memory** and remember alternatives so search can backtrack. **The solution is a path to a goal.**
- Path may be irrelevant if only the final configuration is needed
 - 8-queens
 - IC design
 - Network optimization
 - Factory floor layout
 - Job shop scheduling
 - Automatic programming
 - Cop planning etc.

Local Search


- Use a single current state and move only to neighbours.
 - No track of the paths
 - No track of the set of states that have been reached
- Not systematic → Might never explore a portion of the search space where the solution exists.
- Advantages:
 - Use little space
 - Can find reasonable solutions in large or infinite (continuous) state spaces for which the other algorithms are unsuitable.

Optimization

- Local search is often suitable for optimization problems. Search for the best state by optimizing an **objective function**.
- **$F(x)$ where often x is a vector of continuous or discrete values.**
- Begin with a complete configuration.
- **A successor of state S is S with a single element changed.**
- **Move from the current state to a successor state.**
- Low memory requirements, because the search tree or graph is not maintained in memory (paths are not saved).

Examples – 8 Queens

- Find an arrangement of 8 queens on a chess board such that no two queens are attacking each other (A queen attacks any piece in the same row, column, or diagonal.).
- Start with some arrangement of the queens, one per column.
- $X[j]$: row of queen in column j
- Successors of a state: move one queen
- $F(x)$: # pairs attacking each other



18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	13	16	13	16	16
14	17	15	14	16	16	16	16
17	16	18	15	14	15	16	16
18	14	15	15	14	16	16	16
14	14	13	17	12	14	12	18

Examples – Traveling Salesman Problem

- Visit each city exactly once.
- Start with some ordering of the cities.
- State representation – order of the cities visited.
- Successor state: a change to the current ordering
- $F(x)$: length of the route

Comparison to Tree Search Framework

- Chapter 3: start state is typically not a complete configuration.

Chapter 4: all states are

- Chapter 3: binary goal test;

Chapter 4: no binary goal test, unless one can define one in terms of the objective function for the problem (e.g., no attacking pairs in 8-queens).

- Heuristic function: an estimate of the distance to the nearest goal

Objective function: preference/quality measure – how good is this state?


- Chapter 3: saving paths

Chapter 4: start with a complete configuration and make modifications to improve it

Visualization

- States are laid out in a landscape.
- Elevation corresponds to the objective function value.
- Move around the landscape to find the highest (or lowest) peak.
- Only keep track of the current states and immediate neighbors.

A one-dimensional state-space landscape in which elevation corresponds to the objective function. The aim is to find the global maximum.



Local Search Algorithms

- Strategies for choosing the state to visit next
 - Hill climbing
 - Simulated annealing
- Then, an extension to multiple current states:
 - Genetic algorithms

Hill Climbing (Greedy Local Search)

- Generate nearby successor states to the current state.
- Pick the best and replace the current state with that one.
- Loop

function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

loop do


neighbor \leftarrow a highest-valued successor of *current*

if *neighbor*.VALUE \leq *current*.VALUE **then return** *current*.STATE

current \leftarrow *neighbor*

Hill Climbing (Greedy Local Search)



- Generate nearby successor states to the current state.
- Pick the best and replace the current state with that one.
- Loop



<https://thumbs.gfycat.com/AlienatedImmaculateFlyingsquirrel-mobile.mp4>


Hill Climbing Search Problems

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♔	13	16	13	16
♔	14	17	15	♔	14	16	16
17	♔	16	18	15	♔	15	♔
18	14	♔	15	15	14	♔	16
14	14	13	17	12	14	12	18



Hill Climbing Search Problems



- Here, we assume maximization rather than minimization.
- **Local maximum:** A local maximum is a peak that is higher than each of its neighboring states but lower than the global maximum.



Local maximum

Hill Climbing Search Problems

- **Ridges:** It is a region that is higher than its neighbors but itself has a slope.
 - It is a special kind of local maximum.
 - Results in a sequence of local maxima that is very difficult for greedy algorithms to navigate.
- **Plateau:** the evaluation function is flat, resulting in a random walk.



Variants of Hill Climbing

- **Stochastic hill climbing**

- Chooses at random from among the uphill moves; the probability of selection can vary with the steepness of the uphill move.
- This usually converges more slowly than steepest ascent, but in some state landscapes, it finds better solutions.

- **First-choice hill climbing**

- Implements stochastic hill climbing by generating successors randomly until one is generated that is better than the current state.
- This is a good strategy when a state has many (e.g., thousands) of successors.

Variants of Hill Climbing

- **Random-restart hill climbing**

- Start different hill-climbing searches from random starting positions stopping when a goal is found.
- Save the best result from any search so far.
- If all states have an equal probability of being generated, it is complete with a probability approaching 1 (a goal state will eventually be generated).
- Finding an optimal solution becomes the question of a sufficient number of restarts.
- Surprisingly effective, if there aren't too many local maxima or plateau.

Simulated Annealing


- A hill-climbing algorithm that never makes “downhill” moves is always vulnerable to getting stuck in a local maximum.
- A purely random walk that moves to a successor state without concern for the value will eventually discover the global maximum but will be extremely inefficient.
- Combine the ideas: add some randomness to hill-climbing to allow the possibility of escape from a local optimum.
- **Simulated annealing** wanders around during the early parts of the search, hopefully toward a good general region of the state space
- Toward the end, the algorithm does a more focused search, making a few bad moves.

Simulated Annealing

- Based on a metallurgical metaphor.
- Annealing: harden metals and glass by heating them to a high temperature and then gradually cooling them.
- Start with a temperature set very high and slowly reduce it.
- Run hill-climbing with the twist that you can occasionally replace the current state with a worse state based on the current temperature and how much worse the new state is.
- At the start, make lots of moves and then gradually slow down.

Simulated Annealing

- Generate a random new neighbor from the current state.
- If it's better take it.
- If it's worse, then take it with some probability proportional to the temperature and the delta between the new and old states.



Simulated Annealing


```
function Simulated-Annealing(start, schedule)
  current  $\leftarrow$  start
  for  $t \leftarrow 1$  to  $\infty$  do
     $T \leftarrow$  schedule[ $t$ ]
    if  $T=0$  then return current
    next  $\leftarrow$  a randomly selected successor of the current
     $\Delta E \leftarrow$  Value[next] - Value[current]
    if  $\Delta E > 0$  then current  $\leftarrow$  next
    else current  $\leftarrow$  next only with probability  $e^{\Delta E/T}$ 
```

Simulated Annealing

- Probability of a move decreases with the amount ΔE by which the evaluation is worsened.
- A second parameter T is also used to determine the probability: high T allows worse moves, T close to zero results in few or no bad moves.
- Schedule input determines the value of T as a function of the completed cycles.

Local Beam Search

- Keep track of k states rather than just one, as in hill climbing
- Begins with k randomly generated states
- At each step, all successors of all k states are generated
- If anyone is a goal, algorithm halts
- Otherwise, selects the best k successors from the complete list, and repeats



Local Beam Search

- Successors can become concentrated in a small part of state space
- Stochastic beam search: choose k successors with the probability proportional to the successor's value.
 - Increase diversity

Genetic Algorithms

- Local beam search, but...
 - A successor state is generated by ***combining two parent states***.
- Start with k randomly generated states (**population**).
- A state is represented as a string over a finite alphabet (often a string of 0s and 1s).
- Evaluation function (**fitness function**). Higher = better.
- Produce the next generation of states by **selection**, **crossover**, and **mutation**.


Genetic Algorithms

- Representation of individuals
 - Classic approach: an individual is a string over a finite alphabet, with each element in the string called a gene
 - Usually, binary instead of AGTC as in real DNA
- Mixing number
- Selection strategy
 - Random
 - Selection probability proportional to fitness
 - Selection is made with replacement to make a very fit individual reproduce several times

Genetic Algorithms

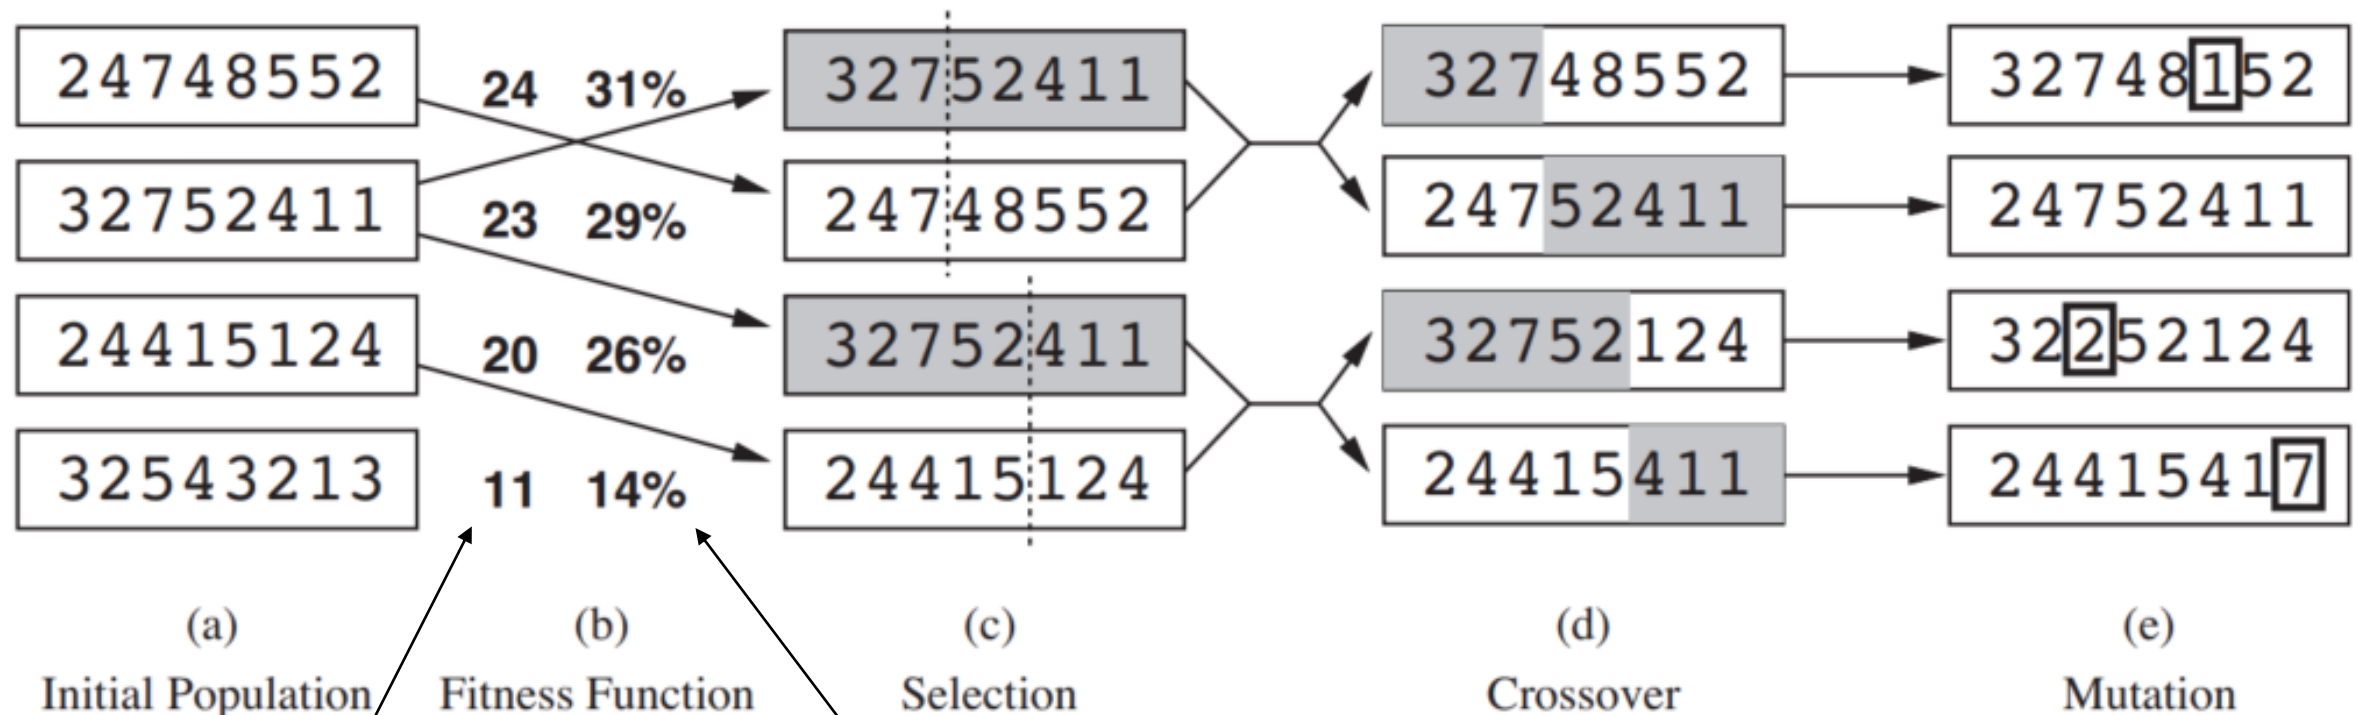
- **Recombination procedure**
 - Random pairing of selected individuals
 - Random selection of **cross-over** points
- **Mutation rate** – How often offspring have random mutations to their representation.
- **The makeup of next generation**
 - Forms with the newly generated offspring
 - **Elitism** - Include a few top-scoring parents from the previous generation in addition to newly formed offspring.
 - **Culling** – Discard individuals below a given threshold.

Example – N-Queens



[16257483]


- Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal



fitness: #non-attacking queens

Probability of being regenerated in next generation based on the fitness value.

Example – N-Queens



- Has the effect of “jumping” to a completely different new part of the search space (quite non-local)

Genetic Algorithms

```
function GA (pop, fitness-fn)
  Repeat
    new-pop = {}
    for i from 1 to size(pop) :
      x = rand-sel(pop, fitness-fn)
      y = rand-sel(pop, fitness-fn)
      child = REPRODUCE(x, y)
      if (small rand prob) : child  $\leftarrow$ 
mutate(child)
      add child to new-pop
    pop = new-pop
until an indiv is fit enough, or out of
time
return best indiv in pop, according to
fitness-fn
```

```
function REPRODUCE(x, y)
  n = len(x)
  c = random num from 1
to n
  return APPEND
(substr(x, 1, c), substr(y,
c+1, n))
```


Comments on Genetic Algorithms

- Genetic Algorithm is a variant of “stochastic beam search”.
- **Positive points**
 - Random exploration can find solutions that local search can't (via crossover primarily).
 - Appealing connection to human evolution.
- **Negative points**
 - Many “tunable” parameters
 - Difficult to replicate performance from one problem to another
 - Lack of good empirical studies comparing to simpler methods
 - Useful on some (small?) problems, but no convincing evidence that GAs are generally better than hill-climbing w/random restarts.